2012

Deep learning using genetic algorithms

Joshua Lamos-Sweeney

Follow this and additional works at: http://scholarworks.rit.edu/theses

Deep Learning Using Genetic Algorithms

Joshua D. Lamos-Sweeney

May 17, 2012

Advisor
Dr. Roger Gaborski

Reader
Dr. Peter Anderson

Observer
Yuheng Wang

Contents

Abstract

Deep Learning networks are a new type of neural network that discovers important object features. These networks determine features without supervision, and are adept at learning high level abstractions about their data sets.

These networks are useful for a variety of tasks, but are difficult to train. This difficulty is compounded when multiple networks are trained in a layered fashion, which results in increased solution complexity as well as increased training time.

This paper examines the use of Genetic Algorithms as a training mechanism for Deep Learning networks, with emphasis on training networks with a large number of layers, each of which is trained independently to reduce the computational burden and increase the overall flexibility of the algorithm.

This paper covers the implementation of a multilayer deep learning network using a genetic algorithm, including tuning the genetic algorithm, as well as results of experiments involving data compression and object classification. This paper aims to show that a genetic algorithm can be used to train a non trivial deep learning network in place of existing methodologies for network training, and that the features extracted can be used for a variety of real world computational problems.

1   Introduction

The creation and training of deep learning networks requires significant computation. This effectively limits the complexity of the networks to be trained. Current solutions for training deep networks are time intensive and limited in the supported neural network architectures[3]. By utilizing a different training method, this paper proposes that more complicated network designs can be attained.

Machine learning is an important field of computing dealing with the acquisition of knowledge. In basic terms, learning is the process of inferring structure in data. Learning is a complex task, which involves making inferences based on a set of data. One major area of research into machine learning involves the creation of feed forward neural networks. Neural networks are conglomerations of small logic units, each of which replicates a simple mathematical function. Feed-forward networks rely on the concept

1

of layers, where all neurons in a layer only output to the next layer. This allows all the data to have a specific, non-infinite chain from input to output. Various methods exist to train these networks to produce a specific output for any specific input. One of the common training methods, known as error propagation, relies on adjusting the network based on how much each neuron contributed to an error, with each neuron passing a part of its error to each neuron which gave it input. By training these networks on a set of data for which the correct output is known, the network will return the appropriate results for similar data.

The first downside of these networks is that they need explicit output design for each category of object to be studied. They also need large data sets to work with, to classify all the differences that can occur. These networks are limited by the simplicity of their network design. Each layer can only perform a simple subset of all possible classifiers, relying on the number of layers to increase the complexity of their classification functions. However, adding layers scales the training time using conventional training algorithms super-linearly. Thus, large, multilayer neural networks cannot be trained in reasonable time. Once trained, the networks are relatively rigid, being able to be applied only to the exact problem domain they are trained against. A network designed to identify three different species of cat in images, for instance, could not be used to find if an image did not contain a cat. A new network, trained on images with and without cats, would have to be constructed.

1.1   Deep Learning

In the simple networks described above, the important information to be learned about the image was given ahead of time. The notations of how each instance of an idea was to be classified was encoded as a set of expected outputs. Labelling the importance of these outputs, and choosing meaningful differences to look for requires human intervention. Features also build on each other. For example, the presence of eyes in an image is a good indicator that there is a person in the image. One of the key points

of deep learning networks is the discovery of features by the algorithm[5]. By discovering relationships in the data set, features can be found more accurately, and by increasing the complexity of the network by adding layers, higher level features, or features concerned less with the structure, and more with the content of the data, can be extracted from the data.

Deep Learning networks extract features by finding common elements in data. By grouping these elements together, a relationship between the elements becomes known. This, in turn, can be connected to another higher level feature. For example, two pixels may always be of similar intensity. These pixels are put together into a feature. As a group of these pixels can share most of their data, one variable can represent the approximate intensity of the group of pixels. If those images were of a face, these pixels might be an eyelid being open or closed. Because of this grouping of low level features, the intensity of the low level pixels, a higher level feature, eyes opened, has been discovered.

Deep Learning networks are currently trained by an algorithm known as Contrastive Divergence. Contrastive Divergence works under the basic principles of conforming the outputs of the network to mimic the training data. By estimating the distance between the current network and an optimum training set, successively closer approximations of a trained algorithm emerge[9]. However, estimating this function requires taking an estimate of the total state of the network. As the network grows more complex, not only does it take longer to train, but each training step increases in computation, limiting the complexity of the networks.

As there is no specified output states, these networks can develop their own rules, requiring less supervision during the learning process than feed forward neural networks. And once trained, these networks can be used more organically, being able to be used on many sub-areas of a problem, including feature extraction, classification, and search.

1.2 Genetic Algorithms

Genetic Algorithms are a type of heuristic search algorithm, based on the concepts of natural selection. The basic operation of a genetic algorithm is simple. A population is created, usually through a random process. The algorithm then runs in a series of steps, known as epochs. Each epoch, new individuals are added to the population, and the worst members of the population are removed. These cycles of survival mimic natural evolution.

There are many variations on Genetic Algorithms, and each one has its own particular lexicon of terms and procedures[4]. This particular algorithm utilizes a constant population, and iterates over a series of generations. Individuals are created using mutation, where a single existing member of the population is subtly changed, and crossover, where two individuals are combined together to create a new individual. Individuals are removed through a selection process, based on the correctness of their solution, known as their fitness.

As a search technique, Genetic Algorithms are heuristics[10]. Heuristics are optimization techniques which are not guaranteed to produce better results than the search methodology they rely on, in this case an exhaustive search, but commonly outperform naive search methods. Heuristics employ assumptions on the structure of the underlying data to attempt to short cut the underlying search methods they are based on. Genetic Algorithms postulate that the search space contains gradients between poor solutions and good solutions. Using this preconception, a Genetic Algorithm can search through the data space focusing on better solutions. If the data followed an ever decreasing error toward a single solution, it would be easy enough to just follow the slope of the decrease towards the optimal value, but even without discontinuities, or locations where the fitness follows unexpected patterns, a simple algorithm can be fooled by local minima.

Because of this, Genetic Algorithms work best on solutions with few discontinuities, but can deal with local minima and other difficulties seen in sub-optimal search problems. The movement of a Genetic Algorithm through a search space can be generalized into two basic motivations, explo-

ration and exploitation. Exploration covers the search space evenly, reducing the chances that the algorithm will stay at a local minima, while exploitation moves the algorithm towards better solutions. With well chosen exploration, the algorithm can be shown to eventually try all possibilities, while well chosen exploitation increases the speed in which the algorithm finds a good solution.

Comparing Genetic Algorithms to the other ways to generate network topologies, there are many advantages to be had[4, 7]. Genetic Algorithms can be trained continuously until a certain condition is made, giving closer approximations to a correct solution. As a search heuristic, a GA will eventually try all combinations. Genetic Algorithms are a good fit for the problem at hand, as there are no discontinuities, but many different good solutions.

Genetically created Deep learning networks can be easily layered, with each layer working on the output of the previous layer, reducing the number of features stored in an image. These networks can be trained sequentially; each network is considered fixed by all networks after it. Using this learning strategy, multi layer networks can scale linearly, as each network can be trained independently.

2   Uses of a Trained Network

2.1   Classification

Once a network is trained against a subset of valid objects, one of the tasks that the network can be used for is classification. In the simplest form, a classification problem can be stated as such: given an object and a set, is this object in the set. One of the underlying principles of the deep learning architecture is the reconstruction of valid objects into their original pattern. Of course, if a random image is sent into such a network, it will not resemble itself very well. However, if the object is a close relation to the objects trained, it should be reconstructed with a high fidelity. Using this, a simple threshold can be established, and if the reconstruction is in error beyond this threshold, it can be declared not in the set.

Figure 1: Reconstruction of a digit and random data.

An example of classification using handwritten data. Two images, one a handwritten zero and the other noise, are compressed and decompressed. The similarity between the input and the output will determine if the image is in the class of handwritten zeroes See Figure 2.1.

2.2 Search

Using conventional networks, search is difficult. Finding an instance of an object in an image is an extremely difficult problem. Many things have to be considered about the object in question, as it may appear significantly different from the training data. Translation, rotation, color aberrations and occlusion are some of the more common problems which stop successful recognition. Also the facing of the object may significantly alter its properties. This requires a significant amount of flexibility in the definition of the object, which is unable to be obtained in simple, single layer feed forward neural networks. This kind of problem requires a large number of layers, to store all the information about the object. Multilayer deep learning networks can be used here to represent objects in the image.

Once a suitable network has been trained on the data, using the network to determine the location of the object or objects is relatively simple. By taking patches of the image, the network can be run on each patch. Patches containing predominantly the object in question will be reconstructed with a high fidelity, in pixels where the object is present. By thresholding areas which return good reconstructed values, a map of the patches containing the

object can be created. The intersection of these patches, naturally, contains the object to be found.

This algorithm is written for image based search, as this paper is focused on image manipulation; However there is no reason for similar techniques cannot be used on other types of data.

2.3   Data Reduction

Deep Learning networks can recreate close approximations of their original objects from a compressed form. Ignoring the cost of the network itself, which is a substantial single cost, this algorithm will perform compression on any object given to it.

If one assumes a perfectly trained matrix, such that each output is recreated as perfectly as possible, it is easy to derive that the algorithm should perform nearly to the optimal level of entropy, as stated in Shannon's Theorem. If there is only one object, it will be perfectly recreated using no data, as the algorithm can simply record all of the data in the bias term. If there are two objects, the two algorithm needs two non-zero rows, which requires one bit of data to store. It can, if perfectly trained, save a 0 in the case for the first object, and a 1 in the case of the second. The bias row contains the first object, while the one reverse row contains the second object minus the first object. Thus reconstruction is perfect. Pushing the limits past Shannon's theorem in this thought experiment cleanly goes to a lossy approximation, where the optimal result is the median point between the two objects, thus both objects have lost half their data. If each training exemplar is given equal weight during training, and all the data to be stored is perfectly trained in the algorithm, the result must be that each is given an equal percentage of the entropy needed to create a perfect reconstruction. Whether this training set can be achieved can be tested while training the network. Choosing a minimal value, v, where v is the ceiling of the entropy of the testing set, as the number of values in the intermediate state given a perfectly trained matrix, will result in the recreation of all testing data perfectly, with a storage cost per training entity equal to the ceiling of the

theoretical minimum amount of data which needs to be stored.

Choosing a minimal value, η, where η is the ceiling of the entropy of the testing set of size θ, as the number of values, or in this case, as the value is a double precision value capped at -1 to 1. Given a perfectly trained matrix of size κ, the size of the data in the testing set, by η, it can restore each input from a storage size of η with no loss of precision, with a total storage cost of $(\kappa \times \eta) + (\eta \times \theta)$. Thus, given a properly trained network, a deep learning network can become a near perfect algorithm for lossless file size reduction, having only $\kappa \times \eta$ overhead over the theoretical minimum storage requirement.

2.4   Data Mining

A deep learning network is designed to create a feature subset of the original object. As training a deep learning network relies on keeping the maximal amount of information about the differences in the patterns of objects in these categories, it can be assumed that the internal representation of the objects keeps as much of the variance between the objects as possible. Thus, this form contains all the variance needed to perform tasks related to the sorting of these objects, such as primary component analysis, and data mining techniques. These patterns have significantly less dimensionality than the input images, allowing for faster, more accurate partitioning of the data space.

Modern Data Mining algorithms use simple dimensionality reduction tools to reduce the size of the data space to be searched while maximizing the distance between points. By reducing the amount of data to compare, these algorithms reduce the run time of comparisons, and reduce confusion, leading to better solutions. By automatically extracting features which are important to recreating the image, a Deep Learning network can perform much the same task, reducing the overall dimensionality of the data, without reducing the importance of the data to the classification of the image.

3 Network Implementation

The deep learning networks created to test these theories can be represented as matrices. Each matrix contains one row for each input, plus a bias row, and one column for each output, plus a bias column. Encoding of data is performed by a simple matrix multiplication, while decoding the data back into the reconstructed input is performed by multiplying the output of the first transformation by the matrix transpose of the input matrix. Both these operations are scaled such that the resulting matrices are in the range of zero to one, though no such limitation is applied to the transformation matrix.

3.1 Basic Network Design

The basic architecture of the networks is a Restricted Boltzmann Machine. An RBM is a network where every neuron in a layer is connected to every neuron in the next successive layer. No other connections are allowed. A neural network in this configuration can be easily represented as a matrix, where rows are the input neurons, and columns are the output neurons. Values in this representation are the weights assigned to the connections.

By themselves, layers in an RBM are not very powerful. However, by increasing the number of layers, complex logic can be represented by the series of operations[3]. Training all the layers simultaneously is infeasible, but if these layers are frozen once computed, the computation time required to compute each layer becomes linear with respect to the number of layers. If each layer is treated as the final layer, reducing the input into a smaller number of features, it can be trained directly on the input. Once it has frozen, its output can be considered the input for the next layer. Once these nested layers have been trained, calculating the final feature set, as well as the reconstructed input, requires applying each matrix in turn, as can be seen in Figure 2. A simple example of a conversion can be seen in Table 1.

Figure 2: An flow chart explaining the order of operations reconstructing an input using a multi-layer network.

Neuron	1	2	3
1	1.0	0.0	0.2
2	-0.5	1.0	-1.0

Table 1: An example matrix. Each entry is a connection between an input(row) and an output(column). Assuming an input vector of [1.0 0.5 1.0], the output is [1.2 -1.0]. Reconstructing the input using the inverse results in a reconstructed input of [1.7 -1.0 1.24].

3.2 Genetics and Network Design

The heart of a Genetic Algorithm is the fitness function, the function which determines how close an individual is to the goal. In this case, the primary driver is the error in the reconstruction of the input. What is important is the noticeable total error, with small errors across the image more acceptable than sections being far off. From this, the choice of mean square error between the input and the reconstruction was obvious, giving the fitness algorithm described in Algorithm 1.

Given a population of matrices, mutation and crossover can be performed as relatively simple operations. Mutation is performed by selecting a single element from the list and replacing it with a new element(Algorithm 4). Crossover was performed as two point crossover. Conceptually, the matrix was linearised, and two points were chosen randomly in the matrix. Two children were created, with all values outside the two points taken from one parent, and all values between the two points taken from the other parent(Algorithm 3).

Algorithm 2 Fitness

Individual individual **Set** images

$\quad \forall$ image $\in$ images

$\quad$ reconstructed $\leftarrow$ individual(image)

$\quad$ fitness $\leftarrow 0$

$\quad$ **for** $\forall$ x,y **do**

$\quad\quad$ meanSquareError $= ($individual$_{(x,y)}-$reconstructed$_{(x,y)})^2$

$\quad\quad$ fitness $=$ fitness $+$ meanSquareError

$\quad$ **end for**

$\quad$ **return** fitness

Algorithm 3 Crossover

Individual parent1, parent2

Require: both parents are matrices of size m by n
   child $\leftarrow \emptyset$
   $i \leftarrow randomrange1..(m \times n)$
   $j \leftarrow randomrange1..(m \times n)$

   for $\forall$ parent$_{(x,y)}$ **do**
      if $(x \times m) + y > i$ **and** $(x \times m) + y < j$ **then** child$_{(x,y)} \leftarrow$ parent1$_{(x,y)}$
      else child$_{(x,y)} \leftarrow$ parent2$_{(x,y)}$
      end if
   end for
   return child

Algorithm 4 Mutation

Individual parent

Require: parent is a matrix of size m by n
   child $\leftarrow$ parent
   $i \leftarrow randomrange1..(m)$
   $j \leftarrow randomrange1..(n)$
   $k \leftarrow randomgaussian$
   child$_{(i,j)} \leftarrow k$
   return child

3.3   Sparse Network Design

Once testing had begun, a pattern in the successful children emerged. Most of the values in the matrices were very close to zero. This becomes apparent looking at the degenerate case where the difference in size between the input and the output matrix is 0. The optimal matrix in this case is the identity matrix, which contains only one non-zero entry per row.

Using this assumption, the genetic sequence of an individual can be reworked. Instead of generating and keeping an entire matrix, a list of non-zero elements can be kept and computed. This saves both space, in reducing the amount of data needed to store the matrix, and the time complexity of computing the matrix multiplication.

In the new model, each non-zero element is a gene. Instead of a fixed size genotype, a variable sized genotype, made up of the non-zero elements in the matrix, is used. Each gene contains a row and column index as to where it would be in the matrix, and a magnitude. The initial population is composed of individuals with one randomly chosen gene. Crossover can increase the length of the genotype, while mutation does not change the genotype length. A penalty has to be added to the fitness to discourage trivial additions to the genotype. This penalty must remain relatively small to not discourage creation of more correct networks, but large enough to remove extraneous elements. The updated fitness calculation can be seen in Algorithm 5.

Though this makes computation faster, it complicates mutation and crossover. The first important change is in the varying of the size of the matrices. Adding a new element during mutation would be a relatively large change, changing both the number of non-zero elements and adding a new element chosen at random. Because of this, altering the size of the gene list falls to the crossover operator. Meanwhile the mutation operator needs to mutate more intelligently.

Mutation using the sparse network is a two step process. The child is created first by removing a gene from the parent, and then adding a new gene, with randomly chosen position and magnitude, to the child. This can

be seen in Algorithm 7.

Crossover needs to be able to create both larger and smaller children, but without biasing the solutions towards larger or smaller solutions. This problem is solved by giving each gene a 50% chance of being passed on to the child. On average, the child will have the average number of genes of the parents, but can in rare cases have as many genes as both parents combined. If both parents attempt to give a gene in the same location to a child, a coin toss is held to determine which copy of the gene is given to the child. By adding half the genes of each parent, the child, statistically, should maintain approximately the size of the parent, while allowing significant variation in length. The pseudocode for the variation can be seen in Algorithm 8.

Algorithm 6 Sparse Fitness

Individual individual **Set** images **Number** sparseness limit

   $\forall$ image $\in$ images
   reconstructed $\leftarrow$ individual(image)
   fitness $\leftarrow 0$
   for $\forall$ x,y **do**
      meanSquareError $= (\text{individual}_{(x,y)} - \text{reconstructed}_{(x,y)})^2$
      fitness $=$ fitness $+$ meanSquareError
   end forsparsenessRate $=$ element count / sparseness
      sparsenessPenalty $= 2^{sparsenessRate}$
      fitness $=$ fitness $\times$ sparsenessPenalty
   return fitness

Algorithm 7 Sparse Mutation

Individual parent

   child $\leftarrow$ parent

   select random $x \in$ child
   create gene $\xi \leftarrow random$
   replace child$_x$ with ξ
   return child

Algorithm 8 Sparse Crossover
***

Individual parent1, parent2

    child $\leftarrow \emptyset$
    $\forall$ parent1$_x$
    $\xi \leftarrow random$
    if $\xi <$ parent1$_{prob}$ **then** child$_x \leftarrow$ parent1$_x$
    end if

    $\forall$ parent2$_x$
    $\xi \leftarrow random$
    if $\xi <$ parent2$_{prob}$ **then**
        if parent1$_x \notin$ child **then** child$_x \leftarrow$ parent2$_x$
        else$\xi \leftarrow random$
            if $\xi > 0.5$ **then** child$_x \leftarrow$ parent2$_x$
            end if
        end if
    end if
    return child
***

4 Results

The algorithm was trained and tested on greyscale image data, with each value being normalized in the floating point range 0.0 to 1.0. This data, as a set of vectors, was used as the input to the algorithm. The reconstruction of the algorithm was used to determine the fitness of the individuals. See Section 3.2 and Algorithm 1 for a discussion of the use and calculation of fitness.

4.1 Parameters

During testing of the data, many ideas were used to attempt to minimise the time used by the algorithm.

4.1.1 Mutation

An early attempt to increase the convergence rate of the algorithm was to increase the number of genes mutated with each mutation. This was successful in the short term, giving faster reduction in the first few moments. However, these gains were not sustainable, and the algorithm quickly slowed, ending with a significantly worse reproduction of the image in any non-trivial test. The early gains suggest that a heuristic approach, e.g. Simulated Annealing, to varying mutation rates may increase performance.

4.1.2 Elitism

When using a Genetic Algorithm, ensuring high fitness values among the population is important. However, if a large group of individuals become very similar. As these similar individuals replace members of the population, the diversity of the population decreases. This can be mitigated by removing old members of the population. Removing elitism, the ability of very old individuals to compete with a new generation, can work to reduce the delay of local minima, however testing in this problem space did not show any improvement.

4.1.3 Enforced Sparseness

As this algorithm has a variable length genetic code, and the length of a member of the population changes the computational complexity of the fitness calculation, it is reasonable to conclude that a barrier on the complexity of the length of the genotype would increase the accuracy of the solution, to a point. However, simply adding a maximum percentage non-sparseness limits solutions which may be better using more than the maximum number of non-zero elements.

A compromise, where the individual was penalized for using non-zero elements was tested, where each gene gave an ever increasing penalty to the fitness of the solution. This penalty was multiplied by the error in the solution, thus making even small improvements capable of reducing the overall fitness value, even in individuals with large numbers of genes.

When the penalty value was scaled well, this resulted in solutions which converged in approximately the same number of generations, but significantly faster in total computation time.

Results dropped off somewhat when the penalty was very small, but kept a noticeable improvement in the overall computation speed. However, if the penalty was large, the computation would fail to find a good solution.

4.1.4 Multithreading

The bulk of the application's running time was spent in the generation of fitness values. Placing each fitness calculation into a thread, and calculating the values simultaneously, resulted in a significant increase in the speed of calculations on multi-threaded systems, with only a minor change to the underlying code. On smaller examples, this speed-up was nearly linear, but as the size of the inputs and resulting matrices increased, the bottleneck moved to memory on the test machines.

4.1.5 Selection Method

The selection of individuals to keep and to remove from a population is an important factor to the design of a genetic algorithm. The simplest choice,

Figure 3: Comparison of Selection Criterion, based on time and reconstruction error.

selecting the most fit elements in each generation, can fall victim to losing variation in the population, leading to slow computation. As mentioned in section 1.2, a loss of variation in the population can reduce the effectiveness of the algorithm at escaping local minima. Tests done on the face images, however, did not show any indication of lack of variation, and outperformed their counterparts using tournament and proportional selection strategies, which can be seen in Figure 3.

18

4.1.6   Layers

Each level of data reduction can be calculated assuming the previous layer is fixed. However the benefit of calculating intermediate steps to the final calculation time and accuracy was unknown. Looking at the accumulated data, as can be seen in Figure 12, the total reconstruction error stays low for multilayer tests. As the mappings in each layer are more straightforward, the GA gains a significant time and accuracy boost by repeatedly taking small, incremental decreases in the compression of the data. Using small steps, the algorithm can compute each layer quickly, while keeping reconstruction fidelity higher than if it had attempted to get to that number of features using less layers. Also, layers can be calculated until the point is reached where important data is being lost, which can be inferred from a sudden increase in error rate, or a decrease in the speed of calculating a layer to an error rate e. This effect is quite noticeable, as can be seen in Figure 4.

4.2   Handwriting

The first item used for training with this algorithm was handwritten numeral characters, using Professor Peter Anderson's digit data. The basic handwriting samples contained digits in a black and white format. Each digit was encoded on a 16x16 black and white image, giving 256 individual pixels for the algorithm.

Previous attempts have been made to determine the efficiency of neural networks, genetic algorithms, and feature selection using this data set[1, 2]. The results of training a matrix to determine the numeral zero were promising. The algorithm reconstructed numbers not in the training set with a high efficiency from a subset of 100 features, as can be seen in Figure 5.

4.3   Face Images

The face images were the next step in complexity. These images contained faces, placed in the center of the image. Unlike the handwriting data, these

Reconstruction Error Per Layer
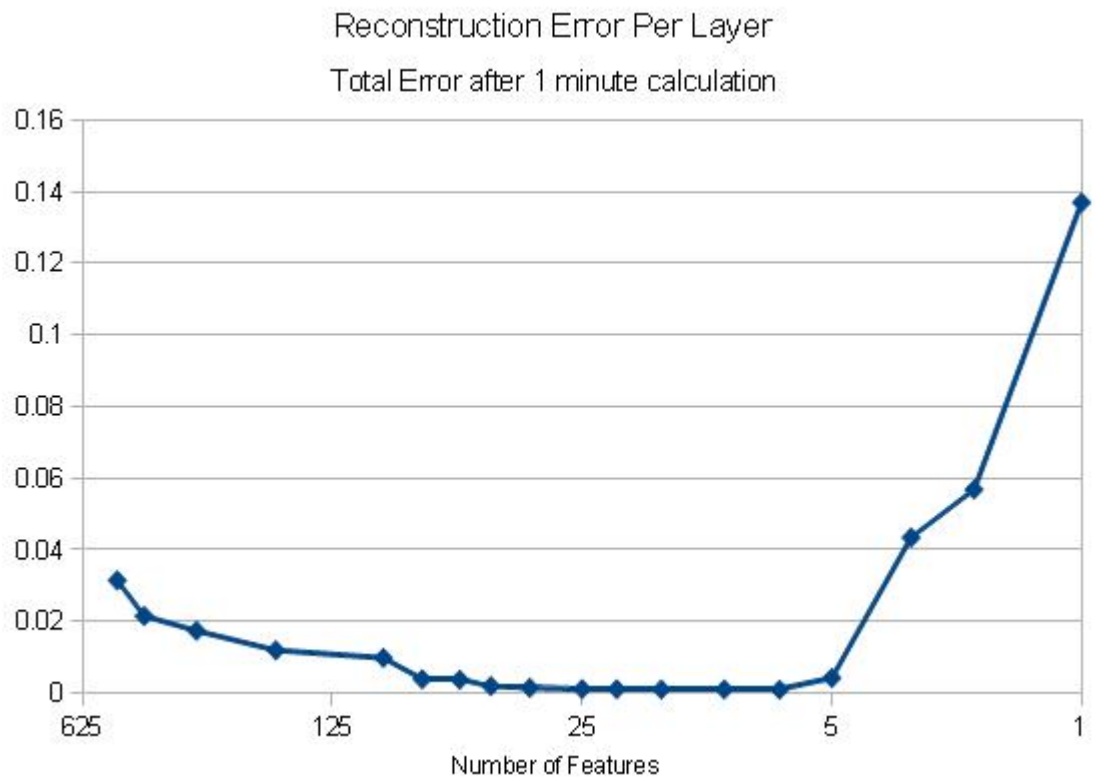
Total Error after 1 minute calculation

Figure 4: Error in reconstruction from the previous layer, from an n feature matrix.

Figure 5: Left to Right: Input, Naive RBM reconstruction, Sparse Matrix reconstruction. Top to Bottom: Training set, Training set, Testing set

20

Figure 6: A set of reconstructed faces from a training image. From left to right, top to bottom the images are the original at `625` elements, followed by `400, 200, 50, 10,` and `1` element.

images were greyscale, and contained significant amounts of features. An example of a reconstructed face image can be seen in Figure 6.

4.4 Cat Images

Another test set examined for determining the attributes of this algorithm was a series of felines. These images were two orders of magnitude larger than the previous images, and contained more background than the previous figures. The complete training set can be seen in Figure 10.

Though the algorithm performed well classifying the part of the set reserved for testing, correctly identifying six of the seven test images, the lack of size of the training set led to difficulty keeping features, reducing the resulting matrix to a combination of the input images, as can be seen in figure 13.

Figure 7: Faces reconstructed from 90 elements. Left is from training set, Right is from testing set.


Figure 8: Reconstruction of 90 elements composed of zeros, random noises, and ones, respectively.


Figure 9: Faces reconstructed from 90 elements with 10 elements replaced. From left to right: Control, 0.0, 1.0, Gaussian Random Distribution

Figure 10: Five cat images used as the training set.

4.5 Large Faces

The final set of images used for testing were a large face database[11]. This database contained significantly more images to train against, and increased in size to 112 by 92, giving the algorithm an input vector over 10,000 entries long. Given 10 different face images, the expected bitwise entropy is 3504 bits, over the total image.

Reducing the feature set quickly introduced noise into the testing set, and though the overall error remained small, the images took on a very grainy appearance when reduced below 7500 elements, as can be seen in Figure 11.

4.6 Missing Data Tests

A simple, visual representation of the effectiveness of the algorithm is to look at how it deals with changing features. Changing blocks of features, if the algorithm is looking at high level features in its reconstruction, should drastically change the output image. These features, when altered, should remove the individuality of objects from the bulk of the learned class. Figure 8 shows that simple changes to the entire feature matrix change the outcome of the reconstruction drastically. At 90 features, the data of the face has been reduced to less than 15% of its original total, yet the images generated with non-zero weights do not look very human. Taking away 1/9th of the values (Figure 9) shows interesting results. Even missing over 10% of the data, the resulting images still look similar. However, when noise is added to the feature set in those 10 places, the left cheek and hairline become very noisy. In the first two cases, the results are filled in by expected human

23

Figure 11: Large face images after reduction to 6500 features. Images courtesy of AT&T Labratories Cambridge[11]. Top two images are from the training set.

traits, while in the latter, the noise follows unexpected patterns, reducing the quality of the image. This shows the features extracted by the algorithm are not simply copying blocks of data, but generating underlying connections in the data.

5   Entropy

The results of running this algorithm have shown the theoretical limits of reconstruct-ability. This algorithm, as stated earlier, approaches the limits of compression via information theory. Looking at each variable in the input set as an independent variable x, the amount of data required to reconstruct each x can be found by looking at the probability of a particular value being correct. If there is only one training image, the total cost of this reconstruction is zero. This is obvious, as a training set of one needs no data to reconstruct it perfectly.

If the resulting size of the Deep Learning network is insufficient to contain the data represented in the testing set, training error increases drastically. See Figure 12, an example of a network being trained past the point where the sum of the data could be stored in the compressed form. Once this state is reached, the algorithm quickly loses definition, as can be seen in Figure 2. From this point onwards, the image loses features from least to most important for the overall reconstruction of the image, seamlessly moving from an attempt at lossless to lossy compression (See Section 2.3 for a discussion of compression through feature extraction).

6   Future Work

Due to the unlabelled nature of the deep learning networks, there is no reason that they cannot be created and differentiated automatically. A system which, when given an input image, classifies it against the known categories, and once it has been classified as such, recomputes the network using the new image as training may be able to create object categories on the fly. This can be accomplished with a finite number of networks, or new networks
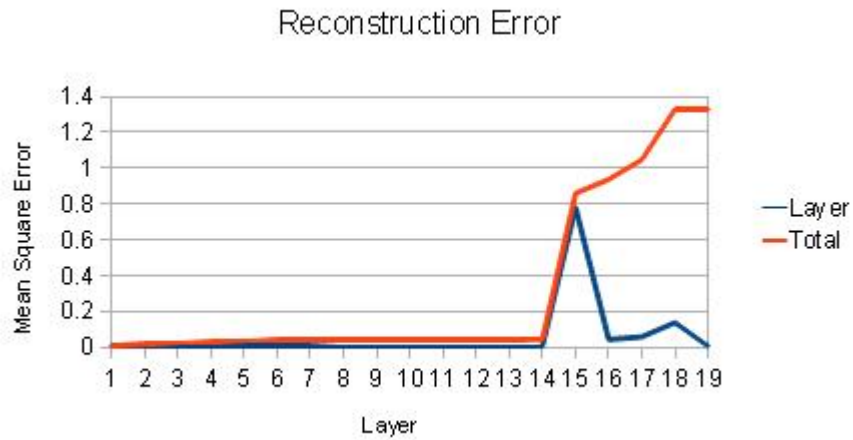
Figure 12: Error rates by layer on a 20 layer run, where each layer reduced the amount of data by approximately 20%. As the amount of data stored drops below a threshold, the reconstruction error per layer increases drastically.
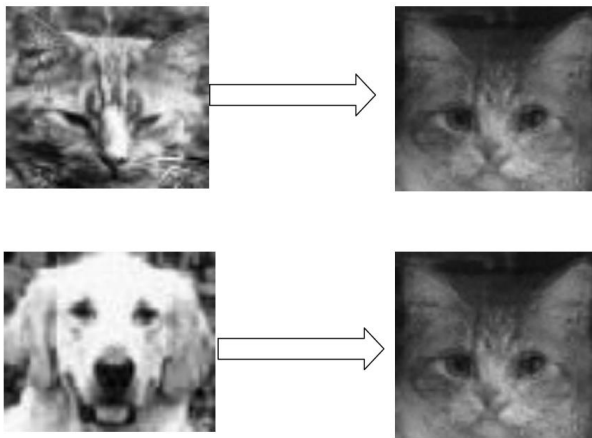

Figure 13: When data is lost, the resulting image will take on features of an amalgam of the training set.

can be created when the classification falls above a certain threshold, thus creating an adaptive network of object classifications, independent of user input. This can allow an algorithm to create and classify new object types on the fly, with the fluidity of more robust data mining applications, while maintaining the speed and flexibility of computation of the deep learning network.

As stated earlier, this algorithm does not take any specific optimizations for the domain of imaging, and can work without change with any domain based on fixed length sets of floating point data. With domain specific optimizations, it may perform significantly better within its domain.

References

[1] Anderson, P., Gaborski, R., Tilley, D., & Asbury, C.(1993). "Genetic Algorithm Selection of Features for Handwritten Character Identification." Artificial Neural Nets and Genetic Algorithms. Proceedings of the International Conference.

[2] Anderson, P. , & Gaborski, R. (1993). The polynomial method augmented by supervised training for hand-printed character recognition. Artificial Neural Nets and Genetic Algorithms. Proceedings of the International Conference, 101-106.

[3] Arel, I. , Rose, D. , & Karnowski, T. (2010). Deep machine learning- a new frontier in artificial intelligence research. Ieee Computational Intelligence Magazine, 5(4), 13-18.

[4] Baeck, T. , Hammel, U. , & Schwefel, H. (1997). Evolutionary computation: Comments on the history and current state. IEEE Transactions on Evolutionary Computation, 1(1), 3-17.

[5] Bengio, Y. (2009). Learning Deep Architectures for AI (Vol. 2, pp. 127). Foundations and Trends in Machine Learning.

[6] Deep Learning. (2006). Retrieved April 15th, 2012, from http://deeplearning.net/

[7] De Jong, K. A. (2006). Evolutionary computation : a unified approach. Cambridge, Mass.: MIT Press.

[8] Hinton, G. (2010). A Practical Guide to Training Restricted Boltzmann Machines (Version 1 ed.). University of Toronto: Department of Computer Science.

[9] Hinton, G. , Osindero, S. , & Teh, Y. (2006). A fast learning algorithm for deep belief nets. Neural Computation, 18(7), 1527-1554

[10] Srinivas, M. , & Patnaik, L. (1994). Genetic algorithms: A survey. Computer, 27(6), 17-26.

[11] The Database of Faces. AT&T Laboratories Cambridge. Retrieved 4/20/2012, 2012, from http://www.cl.cam.ac.uk/research/dtg/attarchive/facedatabase.html